

1. Tahap Pertama

```
import pandas as pd

# Load the dataset
file_path = '/content/insurance.csv'
df = pd.read_csv(file_path)

# Display the first few rows of the dataset and dataset information
df_info = df.info()
df_head = df.head()

df_info, df_head
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
(None,
   age    sex    bmi  children  smoker    region    charges
0    19  female  27.900         0     yes  southwest  16884.92400
1    18   male  33.770         1     no   southeast  1725.55230
2    28   male  33.000         3     no   southeast  4449.46200
3    33   male  22.705         0     no  northwest  21984.47061
4    32   male  28.880         0     no  northwest  3866.85520)
```

2. Tahap Kedua

```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Split the dataset into features (X) and target (y)
X = df.drop(columns='charges')
y = df['charges']

# Define categorical and numerical columns
categorical_features = ['sex', 'smoker', 'region']
numerical_features = ['age', 'bmi', 'children']

# One-hot encode categorical variables and scale numerical variables
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numerical_features),
        ('cat', OneHotEncoder(), categorical_features)
    ])

# Split the data into train and test sets (80:20 split)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create a pipeline that includes preprocessing and linear regression
pipeline = Pipeline(steps=[('preprocessor', preprocessor),
                           ('model', LinearRegression())])

# Train the model
pipeline.fit(X_train, y_train)

# Make predictions
y_pred = pipeline.predict(X_test)

# Evaluate the model
r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)

r2, mse, mae

```

↔ (0.7835929767120722, 33596915.851361476, 4181.194473753652)

3. Tahap Ketiga

```

import matplotlib.pyplot as plt
import seaborn as sns

# Set plot style
sns.set(style='whitegrid')

# Create subplots
fig, axes = plt.subplots(1, 3, figsize=(18, 6))

# Scatter plot for Age vs. Charges
sns.scatterplot(x='age', y='charges', data=df, ax=axes[0], hue='smoker')
axes[0].set_title('Age vs. Charges')
axes[0].set_xlabel('Age')
axes[0].set_ylabel('Charges')

# Scatter plot for BMI vs. Charges
sns.scatterplot(x='bmi', y='charges', data=df, ax=axes[1], hue='smoker')
axes[1].set_title('BMI vs. Charges')
axes[1].set_xlabel('BMI')
axes[1].set_ylabel('Charges')

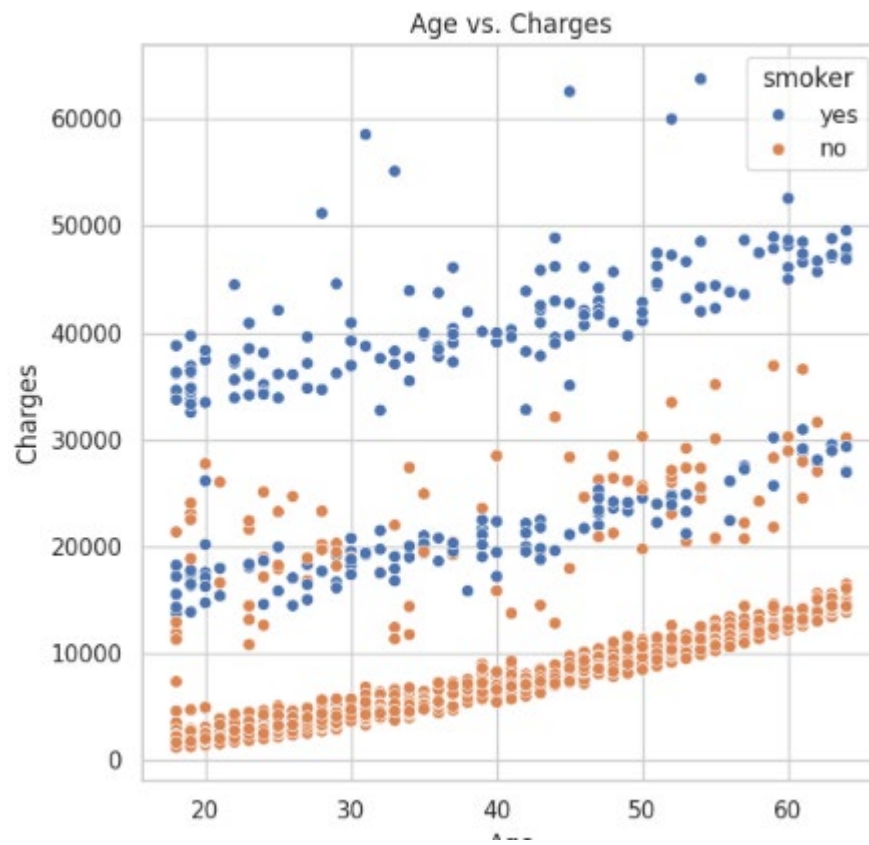
# Bar plot for Smoker vs. Average Charges
sns.barplot(x='smoker', y='charges', data=df, ax=axes[2])
axes[2].set_title('Smoker vs. Average Charges')
axes[2].set_xlabel('Smoker')
axes[2].set_ylabel('Average Charges')

# Show the plots
plt.tight_layout()
plt.show()


```

4. Hasil 1


(4)



5. Hasil 2



6. Hasil 3



7. Tahap 4

```
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

```
# Evaluasi model  
r2 = r2_score(y_test, y_pred)  
mse = mean_squared_error(y_test, y_pred)  
mae = mean_absolute_error(y_test, y_pred)
```

```
# Tampilkan hasil evaluasi  
print(f"R-squared: {r2}")  
print(f"Mean Squared Error (MSE): {mse}")  
print(f"Mean Absolute Error (MAE): {mae}")
```

```
R-squared: 0.7835929767120722  
Mean Squared Error (MSE): 33596915.851361476  
Mean Absolute Error (MAE): 4181.194473753652
```